

## Sample Test | Deep Learning, Spring 2019 | February 21, 2019 | Total Points = 60

All questions have equal points. This is a closed book exam. No electronics or cheat sheets are allowed.

Name: \_\_\_\_\_

1. Assume that the variable "YVALIDATION" has the true labels (1s and 0s) for the Pima Indians Diabetes Dataset and "prediction" variable contains the predictions from a trained neural network model. Given the following code, calculate the precision, accuracy, and recall.

```
1 print(YVALIDATION[0:10])
2
3 prediction = model.predict(XVALIDATION)
4 print('')
5 print(prediction[0:10])
6 print('')
7 print(prediction[0:10].round())
8
9 accuracy = accuracy_score(YVALIDATION[0:10], prediction[0:10].round())
10 precision = precision_score(YVALIDATION[0:10], prediction[0:10].round())
11 recall = recall_score(YVALIDATION[0:10], prediction[0:10].round())
12
13 print("Accuracy: %.2f%%" % (accuracy * 100.0))
14 print("Precision: %.2f%%" % (precision * 100.0))
15 print("Recall: %.2f%%" % (recall * 100.0))
```

```
☞ [1.00 0.00 1.00 0.00 1.00 0.00 1.00 0.00 1.00 1.00]
```

```
[[0.70]
 [0.17]
 [0.70]
 [0.13]
 [0.62]
 [0.23]
 [0.18]
 [0.35]
 [0.62]
 [0.38]]
```

```
[[1.00]
 [0.00]
 [1.00]
 [0.00]
 [1.00]
 [0.00]
 [0.00]
 [0.00]
 [0.00]
 [1.00]
 [0.00]]
```

2. Of the two models (model1 and model2) below, which is for regression and which is for binary classification? Explain.

```
1 model1 = Sequential()
2 model1.add(Dense(12, input_dim=8, activation='relu'))
3 model1.add(Dense(8, activation='sigmoid'))
4 model1.add(Dense(1, activation='linear'))
5
6 model2 = Sequential()
7 model2.add(Dense(12, input_dim=8, activation='sigmoid'))
8 model2.add(Dense(8, activation='relu'))
9 model2.add(Dense(1, activation='sigmoid'))
10
```

3. The code below uses a simple CNN architecture to train a model on the standard MNIST digits dataset (0 to 9 digits).

Complete the code below to have the following architecture:

- (a) 16 convolutional filters in the first layer, each of size 3 x 3, with sigmoid activations
- (b) 4 convolutional filters in the second layer, each of size 3x3, with sigmoid activations
- (c) A dense layer as output layer

```
1 from keras.datasets import mnist
2 from keras.utils import to_categorical
3 from keras import layers, models
4
5 (train_images, train_labels), (validation_images, validation_labels) = mnist.load_
6
7 train_images = train_images.reshape( ( 60000, 28, 28, 1 ) )
8 train_images = train_images.astype( 'float32' ) / 255
9 validation_images = validation_images.reshape( ( 10000, 28, 28, 1 ) )
10 validation_images = validation_images.astype( 'float32' ) / 255
11
12 train_labels = to_categorical( train_labels )
13 validation_labels = to_categorical( validation_labels )
14
15 model = models.Sequential()
16 model.add(layers.Conv2D(                                ))
17 model.add(layers.Conv2D(                                ))
18 model.add(                                              )
19 model.add(layers.Dense(                                ))
20
21 model.compile( optimizer = 'rmsprop', loss = 'categorical_crossentropy', metrics =
22 model.fit( train_images, train_labels, epochs = 1, batch_size = 256 )
```

4. The output below shows the summary of a convolutional neural network constructed.

Answer the following referring to the output:

- (a) How many filters are there in the first convolutional layer?
- (b) How is the total number of parameters in the first layer 160?
- (c) How is the total number of parameters in the second convolutional layer 580?

```
1 from keras.datasets import mnist
2 from keras.utils import to_categorical
3 from keras import layers, models
4
```

```

5 model = models.Sequential()
6 ...
7 model.summary()

```



Layer (type)	Output Shape	Param #
=====		
conv2d_3 (Conv2D)	(None, 26, 26, 16)	160
conv2d_4 (Conv2D)	(None, 24, 24, 4)	580
flatten_2 (Flatten)	(None, 2304)	0
dense_59 (Dense)	(None, 10)	23050
=====		
Total params: 23,790		
Trainable params: 23,790		
Non-trainable params: 0		
=====		

## ▼ 5. What will be the output of the following code?

```

1 import numpy as np
2 x = np.array([[1,2],[3,4]])
3 print np.sum(x)
4 print np.sum(x, axis=0)
5 print np.sum(x, axis=1)

```

## ▼ 6. What will be the output of the following code?

```

1 import numpy as np
2 x = np.array([[1,2,3], [4,5,6], [7,8,9], [10, 11, 12]])
3 v = np.array([1, 0, 1])
4 y = x + v
5 print (x)
6 print (v)
7 print (y)

```



```

[[ 1  2  3]
 [ 4  5  6]
 [ 7  8  9]
 [10 11 12]]
[1 0 1]
[[ 2  2  4]
 [ 5  5  7]
 [ 8  8 10]
 [11 11 13]]

```